

Data Modelling and Databases - Chapter 18 (Book)

- Author: Ruben Schenk
- Date: 26.05.2021
- Contact: ruben.schenk@inf.ethz.ch

18. Concurrency Control

The timing of individual steps of different transactions needs to be regulated in some manner. This regulation is the job of the **scheduler** component of the DBMS, and the general process of assuring that transactions preserve consistency when executing simultaneously is called **concurrency control**.

We begin by studying how to assure that concurrently executing transactions preserve correctness of the database state. The abstract requirement is called **serializability**, and there is an important, stronger condition called **conflict-serializability** that most schedulers actually enforce.

18.1 Serial and Serializable Schedules

18.1.1 Schedules

A **schedule** is a sequence of the important actions taken by one or more transactions. When studying concurrency control, the important read and write actions take place in the main-memory buffers, not the disk.

Example: Let us consider two transactions and the effect on the database the execution of those transactions has. We shall assume that the only consistency constraint on the database state is that $A = B$. Since T_1 adds 100 to both A and B , and T_2 multiplies both A and B by 2, we know that each transaction, run in isolation, will preserve consistency.

The important actions of the transactions T_1 and T_2 are shown in Fig. 18.2:

T_1	T_2
READ(A,t)	READ(A,s)
t := t+100	s := s*2
WRITE(A,t)	WRITE(A,s)
READ(B,t)	READ(B,s)
t := t+100	s := s*2
WRITE(B,t)	WRITE(B,s)

Figure 18.2: Two transactions.

18.1.2 Serial Schedules

A schedule is **serial** if its actions consist of all the actions of one transaction, then all the actions of another transaction, and so on. No mixing of the actions is allowed.

Example: Fig. 18.3 shows one of the two possible serial schedules of the transactions in Fig. 18.2:

T_1	T_2	A	B
		25	25
READ(A,t)			
t := t+100			
WRITE(A,t)		125	
READ(B,t)			
t := t+100			
WRITE(B,t)			125
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		250

Figure 18.3: Serial schedule in which T_1 precedes T_2 .

18.1.3 Serializable Schedules

In general, we say a schedule S is **serializable** if there is a serial schedule S' such that for every initial database state, the effect of S and S' are the same.

Example: Fig. 18.5 below shows a schedule of our previous example that is serializable but not serial. Since all consistent database states have $A = B = c$ for some constant c , it is not hard to deduce that in the schedule of Fig. 18.5, both A and B will be left with the value $2(c + 100)$, and thus consistency is preserved.

T_1	T_2	A	B
		25	25
READ(A,t) t := t+100 WRITE(A,t)		125	
	READ(A,s) s := s*2 WRITE(A,s)	250	
READ(B,t) t := t+100 WRITE(B,t)			125
	READ(B,s) s := s*2 WRITE(B,s)		250

Figure 18.5: A serializable, but not serial, schedule.

On the other hand, consider the schedule of Fig 18.6 below, which is not serializable. The reason we can be sure it is not serializable is that it takes the consistent state $A = B = 25$ and leaves the database in an inconsistent state, where $A = 250$ and $B = 150$.

T_1	T_2	A	B
		25	25
READ(A,t) t := t+100 WRITE(A,t)		125	
	READ(A,s) s := s*2 WRITE(A,s)	250	
	READ(B,s) s := s*2 WRITE(B,s)		50
READ(B,t) t := t+100 WRITE(B,t)			150

Figure 18.6: A nonserializable schedule.

18.1.4 The Effect of Transaction Semantics

To simplify the job of the scheduler, it is conventional to assume that:

- Any database element A that a transaction T writes is given a value that depends on the database state in such a way that no arithmetic coincidences occur.

18.1.5 A Notation for Transactions and Schedules

If we assume "no coincidences", then only the reads and writes performed by the transaction matter, not the actual values involved. Thus, we shall represent transactions and schedules by a shorthand notation, in which the actions are $r_T(X)$ and $w_T(X)$, meaning that transaction T reads, or respectively writes, database element X . Moreover, since we shall usually name our transactions T_1, T_2, \dots we adopt the convention that $r_i(X)$ and $w_i(X)$ are synonyms for $r_{T_i}(X)$ and $w_{T_i}(X)$, respectively.

Example: The transactions of Fig. 18.2 can be written as:

$$\begin{aligned} T_1 &: r_1(A); w_1(A); r_1(B); w_1(B); \\ T_2 &: r_2(A); w_2(A); r_2(B); w_2(B); \end{aligned}$$

To make the notation precise:

- An **action** is an expression of the form $r_i(X)$ or $w_i(X)$, meaning that transaction T_i reads or writes, respectively, the database element X .
- A **transaction** T_i is a sequence of actions with subscript i .
- A **schedule** S of a set of transactions \mathcal{T} is a sequence of actions, in which for each transaction T_i in \mathcal{T} , the actions of T_i appear in S in the same order that they appear in the definition of T_i itself. We say that S is an **interleaving** of the actions of the transactions of which it is composed.

18.2 Conflict-Serializability

Schedulers in commercial systems generally enforce a condition, called "**conflict-serializability**". It is based on the idea of a **conflict**: a pair of consecutive actions in a schedule such that, if their order is interchanged, then the behavior of at least one of the transactions involved can change.

18.2.1 Conflicts

In what follows, we assume that T_i and T_j are different transactions, i.e., $i \neq j$:

- $r_i(X); r_j(Y)$ is never a conflict, even if $X = Y$, since neither of these steps change the value of any database element.
- $r_i(X); w_j(Y)$ is not a conflict provided $X \neq Y$, since if T_j should write Y before T_i reads X , the value of X is not changed. Also, the read of X by T_i has no effect on the value T_j writes for Y .
- $w_i(X); r_j(Y)$ is not a conflict if $X \neq Y$, for the same reason as (2).
- Similarly, $w_i(X); w_j(Y)$ is not a conflict as long as $X \neq Y$.

On the other hand, there are three situations where we may not swap the order of actions:

1. Two actions of the same transaction, e.g., $r_i(X); w_i(Y)$, always conflict. The reason is that the order of actions of a single transaction are fixed and may not be recorded.
2. Two writes of the same database element by different transactions conflict. That is, $w_i(X); w_j(X)$ is a conflict. If we swap the order, then we leave X with the value computed by T_i instead of T_j as before.
3. A read and a write of the same database element by different transactions also conflict. That is, $r_i(X); w_j(X)$ is a conflict, and so is $w_i(X); r_j(X)$.

We say that two schedules are **conflict-equivalent** if they can be turned one into the other by a sequence of nonconflicting swaps of adjacent actions. We shall call a schedule **conflict-serializable** if it is conflict-equivalent to a serial schedule. Note that conflict-serializable is a sufficient condition for serializability.

Example: Consider the schedule

$$r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B);$$

from our previous example. We claim this schedule is conflict-serializable. Fig. 18.8 below shows the sequence of swaps in which this schedule is converted to the serial schedule (T_1, T_2) , where all of T_1 's actions precede all those of T_2 .

$$\begin{aligned} & r_1(A); w_1(A); r_2(A); \underline{w_2(A)}; \underline{r_1(B)}; w_1(B); r_2(B); w_2(B); \\ & r_1(A); w_1(A); \underline{r_2(A)}; \underline{r_1(B)}; \underline{w_2(A)}; w_1(B); r_2(B); w_2(B); \\ & r_1(A); w_1(A); \underline{r_1(B)}; \underline{r_2(A)}; \underline{w_2(A)}; \underline{w_1(B)}; r_2(B); w_2(B); \\ & r_1(A); w_1(A); r_1(B); \underline{r_2(A)}; \underline{w_1(B)}; \underline{w_2(A)}; r_2(B); w_2(B); \\ & r_1(A); w_1(A); r_1(B); w_1(B); \underline{r_2(A)}; \underline{w_2(A)}; r_2(B); w_2(B); \end{aligned}$$

Figure 18.8: Converting a conflict-serializable schedule to a serial schedule by swaps of adjacent actions.

18.2.2 Precedence Graphs and a Test for Conflict-Serializability

Given a schedule S , involving transactions T_1 and T_2 , perhaps among other transactions, we say that T_1 **takes precedence over** T_2 , written $T_1 <_S T_2$, if there are actions A_1 of T_1 and A_2 of T_2 , such that:

1. A_1 is ahead of A_2 in S ,
2. Both A_1 and A_2 involve the same database element, and
3. At least one of A_1 and A_2 is a write action.

We can summarize these precedences in a **precedence graph**. The nodes of the precedence graph are the transactions of a schedule S . When the transactions are T_i for various i , we shall label the node for T_i by only the integer i . There is an arc from node i to j if $T_i <_S T_j$.

Example: The following schedule S involves three transactions, T_1 , T_2 , and T_3 :

$$S : r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B);$$

If we look at the actions involving A , we find several reasons why $T_2 <_S T_3$. For example, $r_2(A)$ comes ahead of $w_3(A)$. Similarly, if we look at the actions involving B , we find that there are several reasons why $T_1 <_S T_2$. For instance, the action $r_1(B)$ comes before $w_2(B)$.

We therefore end up with the following precedence graph:




Figure 18.9: The precedence graph for the schedule S of the example above.

To tell whether a schedule S is conflict-serializable, construct the precedence graph for S and ask if there are any cycles. If so, then S is not conflict-serializable. But if the graph is acyclic, then S is conflict-serializable, and moreover, any topological order of the nodes is conflict-equivalent serial order.

18.2.3 Why the Precedence-Graph Test Works

Left out.

18.3 Enforcing Serializability by Locks

In this section, we introduce the concept of locking with a simple locking scheme. In this scheme, there is only one kind of lock, which transactions must obtain on a database element if they want to perform any operation whatsoever on that element.

18.3.1 Locks

A locking scheduler, like most types of scheduler, enforces conflict-serializability, which as we learned is a more stringent condition than correctness, or even serializability.

When a scheduler uses **locks**, transactions must request and release locks, in addition to reading and writing database elements. The use of locks must be proper in two senses:

- **Consistency of Transactions** : Actions and locks must relate in the expected ways:
 1. A transaction can only read or write an element if it previously was granted a lock on that element and hasn't yet released the lock.
 2. If a transaction locks an element, it must later unlock that element.
- **Legality of Schedules** : Locks must have their intended meaning: no two transactions may have locked the same element without having first released the lock.

We shall extend our notation for actions to include locking and unlocking actions:

- $l_i(X)$: Transaction T_i requests a lock on database element X .
- $u_i(X)$: Transaction T_i releases its lock on database element X .

Thus, the **consistency** condition for transactions can be stated as: "Whenever a transaction T_i has an action $r_i(X)$ or $w_i(X)$, then there is a previous action $l_i(X)$ with no intervening action $u_i(X)$, and there is a subsequent $u_i(X)$."

The **legality of schedules** is stated as: "If there are actions $l_i(X)$ followed by $l_j(X)$ in a schedule, then somewhere between these actions there must be an action $u_i(X)$."

Example: Let us consider the two transactions T_1 and T_2 that were introduced at the beginning of this chapter. Then Fig. 18.12 below shows one legal schedule of these two transactions:

T_1	T_2	A	B
		25	25
$l_1(A); r_1(A);$ $A := A+100;$ $w_1(A); u_1(A);$		125	
	$l_2(A); r_2(A);$ $A := A*2;$ $w_2(A); u_2(A);$	250	
	$l_2(B); r_2(B);$ $B := B*2;$ $w_2(B); u_2(B);$		50
$l_1(B); r_1(B);$ $B := B+100;$ $w_1(B); u_1(B);$			150

Figure 18.12: A legal schedule of consistent transactions. Unfortunately it is not serializable.

18.3.2 The Locking Scheduler

It is the job of the a scheduler based on locking to grant requests if and only if the request will result in a legal schedule. If a request is not granted, the requesting transaction is delayed. It waits until the scheduler grants its request at a later time. To aid its decisions, the scheduler has a **lock table** that tells, for every database element, the transaction that currently holds a lock on that element.

18.3.3 Two-Phase Locking

There is a surprising condition, called **two-phase locking (2PL)** under which we can guarantee that a legal schedule of consistent transactions is conflict-serializable:

- In every transaction, all lock actions precede all unlock actions.

The "two phases" referred to by 2PL are thus the first phase, where locks are obtained, and the second phase, where locks are relinquished. A transaction that obeys the 2PL condition is said to be a **two-phase-locked transaction**, or 2PL transaction.

T_1	T_2	A	B
		25	25
$l_1(A); r_1(A);$ $A := A+100;$ $w_1(A); l_1(B); u_1(A);$		125	
	$l_2(A); r_2(A);$ $A := A*2;$ $w_2(A);$ $l_2(B)$ Denied	250	
$r_1(B); B := B+100;$ $w_1(B); u_1(B);$			125
	$l_2(B); u_2(A); r_2(B);$ $B := B*2;$ $w_2(B); u_2(B);$		250

Figure 18.13: The locking scheduler delays requests that would result in an illegal schedule.

18.3.4 Why Two-Phase Locking Works

Left out.

18.4 Locking Systems With Several Lock Modes

The locking scheme of Section 18.3 illustrates the important ideas behind locking, but it is too simple to be a practical scheme. The main problem is that a transaction T must take a lock on a database element X even if it only wants to read X and not write it.

We are thus motivated to introduce the most common locking scheme, where there are two different kinds of locks, one for reading (called a **shared lock**), and one for writing (called an **exclusive lock**).

18.4.1 Shared and Exclusive Locks

Let us consider a locking scheduler that uses two different kinds of locks: **shared locks** and **exclusive locks**. For any database element X there can be either one exclusive lock on X , or no exclusive locks but any number of shared locks. If we want to write X , we need to have an exclusive lock on X , but if we wish only to read X we may have either a shared or exclusive lock on X .

We shall use $sl_i(X)$ to mean "transaction T_i requests a shared lock on database element X " and $xl_i(X)$ for " T_i requests an exclusive lock on X ". We continue to use $u_i(X)$ to mean that T_i unlocks X .

The three kinds of requirements - consistency and 2PL for transactions, and legality of schedules - each have their counterpart for a shared/exclusive lock system. We summarize these requirements here:

1. **Consistency of transactions**: A transaction may not write without holding an exclusive lock, and you may not read without holding some lock. More precisely, in any transaction T_i ,
 1. A read action $r_i(X)$ must be preceded by $sl_i(X)$ or $xl_i(X)$, with no intervening $u_i(X)$.
 2. A write action $w_i(X)$ must be preceded by $xl_i(X)$, with no intervening $u_i(X)$.

All locks must be followed by an unlock of the same element.

2. **Two-phase locking of transactions**: Locking must precede unlocking. To be more precise, in any two-phase locked transaction T_i , no action $sl_i(X)$ or $xl_i(X)$ can be preceded by an action $u_i(Y)$, for any Y .
3. **Legality of schedules**: An element may either be locked exclusively by one transaction or by several in shared mode, but not both. More precisely:
 1. If $xl_i(X)$ appears in a schedule, then there cannot be a following $xl_j(X)$ or $sl_j(X)$, for some j other than i , without an intervening $u_i(X)$.
 2. If $sl_i(X)$ appears in a schedule, then there cannot be a following $xl_j(X)$, for $j \neq i$, without an intervening $u_i(X)$.

Note that we do allow one transaction to request and hold both shared and exclusive locks on the same element, provided its doing so does not conflict with the locks of other transactions.

18.4.2 Compatibility Matrices

A **compatibility matrix** is a convenient way to describe lock-management policies. It has a row and column for each lock mode.

The rule for using a compatibility matrix for lock-granting decisions is:

- We can grant the lock on X in mode C if and only if for every row R such that there is already a lock on X in mode R by some other transaction, then there is a "Yes" in column C .

Example: Fig. 18.14 is the compatibility matrix for shared (S) and exclusive (X) locks. The column for S says that we can grant a shared lock on an element if the only locks held on that element currently are shared locks. The column for X says that we can grant an exclusive lock only if there are no other locks held currently.

		Lock requested	
		S	X
Lock held in mode	S	Yes	No
	X	No	No

Figure 18.16: The compatibility matrix for shared and exclusive locks.

18.4.3 Upgrading Locks

Left out.

18.4.4 Update Locks

Left out.

18.4.5 Increment Locks

Left out.

18.5 An Architecture for a Locking Scheduler

Left out.

18.6 Hierarchies of Database Elements

18.6.1 Locks With Multiple Granularity

Recall that the term "database element" was purposely left undefined, because different systems use different sizes of database elements to lock, such as tuples, pages or blocks, and relations.

18.6.2 Warning Locks

The solution to the problem of managing locks at different granularities involves a new kind of lock called a "warning". These locks are useful when the database elements form a nested or hierarchical structure. There, we see three levels of database elements:

1. Relations are the largest lockable elements.
2. Each relation is composed of one or more block or pages, on which its tuples are stored.
3. Each block contains one or more tuples.

The rules for managing locks on a hierarchy of database elements constitute the **warning protocol**, which involves both "ordinary" locks and "warning" locks. We shall describe the lock scheme where the ordinary locks are S and X (shared and exclusive). The warning locks will be denoted by prefix I (for "intention to") to the ordinary locks, for example IS represents the intention to obtain a shared lock on a subelement. The rules of the warning protocol are:

1. To place an ordinary S or X lock on any element, we must begin at the root of the hierarchy.
2. If we are at the element that we want to lock, we need look no further. We request an S or X lock on that element. If the element we wish to lock is further down the hierarchy, then we place a warning at this node. That is, if we want to get a shared lock on a subelement we request an IS lock at this node, and if we want an exclusive lock on a subelement, we request an IX lock on this node. When the lock on the current node is granted, we proceed to the appropriate child. We then repeat step (2) or step (3), as appropriate, until we reach the desired node.

In order to decide whether or not one of these locks can be granted, we use the compatibility matrix of Fig. 18.28 below:

	IS	IX	S	X
IS	Yes	Yes	Yes	No
IX	Yes	Yes	No	No
S	Yes	No	Yes	No
X	No	No	No	No

Figure 18.28: Compatibility matrix for shared, exclusive, and intention locks.

18.6.3 Phantoms and Handling Insertions Correctly

When transactions create new subelements of a lockable element, there are some opportunities to go wrong. The problem is that we can only lock existing items. There is no way to lock database elements that do not exist but might later be inserted.

A **phantom tuple** is one that should have been locked but wasn't, because it didn't exist at the time the locks were taken. There is, however, a simple way to avoid the occurrence of phantoms. We must regard the insertion or deletion of a tuple as a write operation on the relation as a whole. Thus, we must obtain an X on the relation before executing a insertion or deletion.

18.7 The Tree Protocol

Like section 18.6, this section deals with data in the form of a tree. However, here, the nodes of the tree do not form a hierarchy based on containment. Rather, database elements are disjoint pieces of data, but the only way to get to a node is through its parents. B-trees are an important example of this sort of data.

18.7.1 Motivation for Tree-Based Locking

Let us consider a B-tree index in a system that treats individual nodes as lockable database elements.

If we use a standard set of lock modes, like shared, exclusive, and update locks, and we use two-phase locking, then concurrent use of the B-tree is almost impossible.

In most situations, we can deduce almost immediately that a B-tree node will not be rewritten, even if the transaction inserts or deletes a tuple.

Thus, as soon as a transaction moves to a child of the root and observes the situation that rules out a rewrite of the root, we would like to release the lock on the root. The same observation applies to the lock on any interior node of the B-tree.

18.7.2 Rules for Access to Tree-Structured Data

The following restrictions on locks from the **tree protocol**. We assume that there is only one kind of lock, represented by lock requests of the form $l_i(X)$.

1. A transaction's first lock may be at any node of the tree.
2. Subsequent locks may be only acquired if the transaction currently has a lock on the parent node.
3. Nodes may be unlocked at any time.
4. A transaction may not relock a node on which it has released a lock, even if it still holds a lock on the node's parent.

18.7.3 Why the Tree Protocol Works

Left out.

18.8 Concurrency Control by Timestamps

Next, we shall consider two methods other than locking that are used in some systems to assure serializability of transactions:

1. **Timestamping** : Assign a "timestamp" to each transaction. Record the timestamps of the transactions that last read and write each database element, and compare these values with the transactions timestamps, to assure that the serial schedule according to the transactions' timestamps is equivalent to the actual schedule of the transactions.
2. **Validation** : Examine timestamps of the transaction and the database elements when a transaction is about to commit. This process is called "validation" of the transaction. The serial schedule that orders transactions according to their validation time must be equivalent to the actual schedule.

Both these approaches are **optimistic**, in the sense that they assume that no unserializable behavior will occur and only fix things up when a violation is apparent.

18.8.1 Timestamps

To use timestamping as a concurrency-control method, the scheduler needs to associate with each transaction T its **timestamp** $TS(T)$, and each database element X two timestamps and an additional bit:

1. $RT(X)$, the **read time** of X , which is the highest timestamp of a transaction that has read X .
2. $WT(X)$, the **write time** of X , which is the highest timestamp of a transaction that has written X .
3. $C(X)$, the **commit bit** for X , which is true if and only if the most recent transaction to write X has already committed.

18.8.2 Physically Unrealizable Behaviors

Another job of the scheduler is to check that whenever a read or write occurs, what happens in real time could have happened if each transaction had executed instantaneously at the moment of its timestamp. If not, we say the behavior is **physically unrealizable**. There are two kinds of problems that can occur:

1. **Read too late** : Transaction T tries to read database element X , but the write time of X indicates that the current value of X was written after T theoretically expired. That is, $TS(T) < WT(X)$.
2. **Write too late** : Transaction T tries to write database element X . However, the read time of X indicates that some other transaction should have read the value written by T , but read some other value instead. That is, $WT(X) < TS(T) < RT(X)$.

18.8.3 Problems With Dirty Data

There is a class of problems that the commit bit is designed to solve. One of these problems is a **dirty read**. There. transaction T reads X , and X was last written by U . The timestamp of U is less than that of T , and the read by T occurs after the write by U in real time, so the event seems to be physically realizable. However, it is possible that after T reads the value of X written by U , transaction U will abort. Perhaps, U encounters an error condition in its own data.

Thus, although there is nothing physically unrealizable about T reading X , it is better to delay T 's read until U commits or aborts. We can tell that U is not committed because the commit bit $C(X)$ will be false.

18.8.4 The Rules for Timestamp-Based Scheduling

We can now summarize the rules that a scheduler using timestamps must follow to make sure that nothing physically unrealizable may occur.

The rules are as follows:

1. Suppose the scheduler receives a request $r_T(X)$
 1. If $TS(T) \geq WT(X)$, the read is physically realizable.
 1. If $C(X)$ is true, grant the request. If $TS(T) > RT(X)$, set $RT(x) := TS(T)$, otherwise do not change $RT(X)$.
 2. If $C(X)$ is false, delay T until $C(X)$ becomes true, or the transaction that wrote X aborts.
 2. If $TS(T) < WT(X)$, the read is physically unrealizable. Rollback T , that is, abort T and restart it with a new, larger timestamp.
2. Suppose the scheduler receives a request $w_T(X)$.
 1. If $TS \geq RT(X)$ and $TS(T) \geq WT(X)$, the write is physically realizable and must be performed.
 1. Write the new value for X ,
 2. Set $WT(X) := TS(T)$, and
 3. Set $C(X) := \text{false}$.
 2. If $TS(T) \geq RT(X)$, but $TS(T) < WT(X)$, then the write is physically realizable, but there is already a later value in X . If $C(X)$ is true, then the previous writer of X is committed, and we simply ignore the write by T . We allow T to proceed and make no change to the database. However, if $C(X)$ is false, then we must delay T as in point 1.1.2.
 3. If $TS(T) < RT(X)$, then the write is physically unrealizable, and T must be rolled back.
3. Suppose the scheduler receives a request to commit T . It must find all the database elements X written by T , and set $C(X) := \text{true}$.
4. Suppose the scheduler receives a request to abort T or decides to rollback T as in 1.2 or 2.3. Then any transaction that was waiting on an element X that T wrote must repeat its attempt to read or write.

18.8.5 Multiversion Timestamps

Left out.

18.8.6 Timestamps Versus Locking

Left out.

18.9 Concurrency Control by Validation

Validation is another type of optimistic concurrency control, where we allow transactions to access data without locks, and at the appropriate time we check that the transaction has behave in a serializable manner. Validation differs from timestamping principally in that the scheduler maintains a record of what active transactions are doing, rather than keeping read and write times for all database elements.

18.9.1 Architecture of a Validation-Based Scheduler

When validation is used as the concurrency-control mechanism, the scheduler must be told for each transaction T the sets of database elements T reads and writes, the **read set**, $RS(T)$, and the **write set**, $WS(T)$, respectively.

Transactions are executed in three phases:

1. **Read** : In the first phase, the transaction reads from the database all the elements in its read set.
2. **Validate** : In the second phase, the scheduler validates the transaction by comparing its read and write sets with those of other transactions.
3. **Write** : In the third phase, the transaction writes to the database its values for the elements in its write set.

To support the decision whether to validate a transaction, the scheduler maintains three sets:

1. **START**, the set of transactions that have started, but not yet completed validation. For each transaction T in this set, the scheduler maintains $START(T)$, the time at which T started.
2. **VAL**, the set of transactions that have been validated but not yet finished the writing of phase 3. For each transaction T in this set, the scheduler maintains both $START(T)$ and $VAL(T)$, the time at which T validated.
3. **FIN**, the set of transactions that have completed phase 3. For these transactions T , the scheduler records $START(T)$, $VAL(T)$, and $FIN(T)$, the time at which T finished. In principle this set grows, but as we shall see, we do not have to remember transactions T if $FIN(T) < START(U)$ for any active transaction U .

18.8.2 The Validation Rules

We may summarize with the following rules for validating a transaction T :

- Check that $RS(T) \cap WS(U) = \emptyset$ for any previously validated U that did not finish before T started, i.e., if $FIN(U) > START(T)$.
- Check that $WS(T) \cap WS(U) = \emptyset$ for any previously validated U that did not finish before T validated, i.e., if $FIN(U) > VAL(T)$.

18.9.3 Comparison of Three Concurrency-Control Mechanisms

The three approaches to serializability that we have considered - locks, timestamps, and validation - each have their advantages. First, they can be compared for their storage utilization:

- **Locks**: Space in the lock table is proportional to the number of database elements locked.
- **Timestamps**: In a naive implementation, space is needed for read- and write-times with every database element, whether or not it is currently accessed.
- **Validation**: Space is used for timestamps and read/write sets for each currently active transaction.

We can also compare methods for their effect on the ability of transactions to complete without delay. The performance of the three methods depends on whether **interaction** among transactions, that is the likelihood that a transaction will access an element that is also being accessed by a concurrent transaction, is high or low:

- **Locking** delays transactions but avoids rollbacks, even when interaction is high. Timestamps and validation do not delay transactions, but can cause them to rollback, which is a more serious form of delay and also wastes resources.
- If interference is low, then neither timestamps nor validation will cause many rollbacks.
- When a rollback is necessary, timestamps catch some problems earlier than validation, which always lets a transaction do all its internal work before considering whether the transaction must rollback.

18.10 Summary of Chapter 18

Consistent Database States

Database states that obey whatever implied or declared constraints the designers intended are called **consistent**.

Consistency of Concurrent Transactions

It is normal for several transactions to have access to a database at the same time. It is the job of the scheduler to assure that concurrently operating transactions also preserve the consistency of the database.

Schedules

Transactions are broken into actions, mainly reading and writing from the database. A sequence of these actions from one or more transactions is called a **schedule**.

Serial Schedules

If transactions execute one at a time, the schedule is said to be **serial**.

Serializable Schedules

A schedule that is equivalent in its effect on the database to some serial schedule is said to be **serializable**.

Conflict-Serializability

A simple-to-test, sufficient condition for serializability is that the schedule can be made serial by a sequence of swaps of adjacent actions without conflicts. Such a schedule is called **conflict-serializable**.

Precedence Graphs

An easy test for conflict-serializability is to construct a **precedence graph** for the schedule. Nodes correspond to transactions, and there is an arc $T \rightarrow U$ if some action T in the schedule conflicts with a later action of U . A schedule is conflict-serializable if and only if the precedence graph is acyclic.

Locking

The most common approach to assuring serializable schedules is to lock database elements before accessing them, and to release the lock after finishing access to the element.

Two-Phase Locking

Locking by itself does not assure serializability. However, **two-phase locking**, in which all transactions first enter a phase where they only acquire locks, and then enter a phase where they only release locks, will guarantee serializability.

Lock Modes

To avoid locking out transactions unnecessarily, systems usually use several lock modes, with different rules for each mode about when a lock can be granted.

Compatibility Matrices

A compatibility matrix is a useful summary of when it is legal to grant a lock in a certain lock mode.

Locking Elements With a Granularity Hierarchy

When both large and small elements - relations, disk blocks, and tuples, perhaps - may need to be locked, a warning system of locks enforces serializability.

Locking Elements Arranged in a Tree

If database elements are only accessed by moving down a tree, as in a B-tree index, then a non-two-phase locking strategy can enforce serializability.