

Data Modelling and Databases - Week 1 (Lecture)

- Author: Ruben Schenk
- Date: 27.04.2021
- Contact: ruben.schenk@inf.ethz.ch

Data in Real World

A **database** is a collection of data, for example information about bank accounts or data on facebook, Amazon, etc.

A **database management software (DBMS)** is software designed to assist in maintaining and utilizing large collections of data.

First Database Example

We have several "whishes" for DBMS:

- *Data Independence*: application should not know how data is stored
- *Declarative Efficient Data Access*: the system should be able to store and retrieve data efficiently, without users worrying about it
- *Transactional Access*: as if there is only a single users using a system that does not fail
- *Generic Abstraction*: Users do not need to worry about all the above issues for each new query

What's the **potential downside** of using a DBMS?

- *Workload mismatch*: maybe your specialized application is not what a certain DBMS is designed for
- *Data model mismatch*: maybe your application cannot be naturally modeled by a given DBMS

Data Model: Relational Model

In this course, we focus on a specific combination - we represent knowledge as a **collection of facts**, and do inference using **mathematical logic**.

Relational Model - Schema

A **database schema** is a set of relation schema, where a **relation schema** is defined by a name and a set of attributes/fields/columns. A **field** or **attribute** is defined by a name and a domain, e.g. Integer, String, etc.

For example:

```
Students(sid:string, name:string, login:string, age:int, gpa:float)
```

The above code defines the relation schema "Students" by the attributes "sid, name, login, age", and "gpa".

Relational Model - Instance

For a relation $R(f_1 : D_1, \dots, f_n : D_n)$, an **instance** I_R is a set of tuples: $I_R \subseteq D_1 \times \dots \times D_n$. Intuitively, an instance is the "content" of a relation if you think about it as a "table". It is important to remember that a relation instance is a **set**, this means we cannot have duplicated tuples and that the order of tuples doesn't matter.

Relational Model - More Concepts

A **candidate key** is the minimal set of fields that identify each tuple uniquely. A **primary key** is one candidate key, marked in a schema by underlining.

Query Language 1: Relational calculus

Union: \cup

$$x \in R_1 \cup R_2 \Leftrightarrow x \in R_1 \vee x \in R_2$$

Difference: $-$

$$x \in R_1 - R_2 \Leftrightarrow x \in R_1 \wedge \neg(x \in R_2)$$

Intersection: \cap

$$R_1 \cap R_2 = R_1 - (R_1 - R_2)$$

Selection: σ

Return tuples which satisfy a given condition c .

$$x \in \sigma_c(R) \Leftrightarrow x \in R \wedge c(x) = \text{True}$$

Projection: $\Pi_{A_1, \dots, A_n}(R)$

Only keep a subset of columns.

Cartesian Product: \times

$$(x, y) \in R_1 \times R_2 \Leftrightarrow x \in R_1 \wedge y \in R_2$$

Renaming: $\rho_{B_1, \dots, B_n}(R)$

Change the name of the attributes of R to B_1, \dots, B_n .

Natural join: \bowtie

$$R_1(A, B) \bowtie R_2(B, C) = \Pi_{A, B, C}(\sigma_{R_1.B=R_2.B}(R_1 \times R_2))$$

If there are **no shared attributes** in a natural join, e.g. $R(A, B, C)$ and $S(D, E)$, then $R \bowtie S = R \times S$.

If two relations **share all attributes**, then $R \bowtie S = R \cap S$.

Theta Join: \bowtie_{θ}

$$R_1 \bowtie_{\theta} R_2 = \sigma_{\theta}(R_1 \times R_2)$$

Equi-Join: $\bowtie_{A=B}$

$$R_1 \bowtie_{A=B} R_2 = \sigma_{A=B}(R_1 \times R_2)$$

It is important to note that relational algebra uses **Bag semantics** instead of set semantics:

- Each relation is a bag of tuples
- You can have duplicated tuples in the same relation
- i.e. set: $\{1, 2, 3\}$, bag: $\{1, 2, 3, 1, 2, 1\}$

It is furthermore important to remember that **bag operator semantics** are different to set operator semantics:

- *Bag Union:* $\{1, 2, 1\} \cup \{1, 2, 3\} = \{1, 1, 1, 2, 2, 3\}$
- *Bag Difference:* $\{1, 2, 1\} - \{1, 2, 3, 3\} = \{1\}$